Anton Ernstsson

Junior Software Developer · Data Engineer · AI Engineer

antonernstsson.com

Stockholm, Sweden / +46 762 541 024 / anton.ernstson@gmail.com
github.com/korv9 / linkedin.com/in/anton-ernstsson

Junior software developer with experience developing and maintaining ETL pipelines from two internships. I run my own data platform where I host my experiments and projects, and I have worked with text clustering both in personal projects and on production data.

Professional experience

Analytics Engineer · Avtalat — Jan 2026 - Jun 2026

LIA internship · same project, continued after the company merger

- Prepared the pipeline for Avtalat's new Freshservice instance, so it works across instances and adapts when fields or tables are added in the source.
- Built Power BI reports and semantic models (DAX, DirectQuery) that replaced manual reporting with self-service.
- Ran a reproducible NLP analysis of incident data, including anonymisation and data quality assessment.

Data Engineer · Fora AB — Nov 2025 - Jan 2026

LIA internship · Freshservice ITSM data in Azure Databricks

- Built a Python and PySpark pipeline that loads 80,000+ records from the Freshservice REST API, covering tickets, problems, changes, agents and departments.
- Modelled the data through Bronze, Silver and Gold into a connected star schema of fact and dimension tables, built for straightforward reporting.
- Scheduled it as a daily Databricks Workflows job with incremental loads, full reloads when needed, validation and error handling.
- Replaced a manual report for the IT service team and gave each department a clearer view of its own tickets and problems than Freshservice offers.

Technical skills

LANGUAGES	Python, SQL, PySpark, Spark SQL, TypeScript
DATA ENGINEERING	Azure Databricks, Delta Lake, dbt Core, DuckDB, Parquet, ETL/ELT, REST APIs, Databricks Workflows, Incremental loads, Star schema, Microsoft Fabric
AI AND BI	NLP, sentence-transformers, UMAP, HDBSCAN, MLflow, pandas, Power BI, DAX
WEB AND DELIVERY	React, Vite, Cloudflare, Git, GitHub Actions, Azure DevOps, pytest, Playwright

Education

AI Developer · JENSEN Yrkeshögskola — 2024 - 2026

Higher Vocational Education (YH) · 400 YH credits · graduated Jun 2026

Python, databases, full-stack development, machine learning and deep learning.

Selected projects

antonernstsson.com · Data portfolio and platform

`Personal project · Main project`

- Built a scalable, deployable site where I host my experiments and projects, and keep adding new ones. Bilingual (Swedish/English), built in React and TypeScript, with projects on Swedish and EU public data such as Swedish politics in numbers, an EU AI Act Observatory and the job market in job ads.
- Built the data platform behind it: Python ingestion that keeps raw files with provenance, dbt models on DuckDB in Bronze, Silver and Gold layers, and JSON and Parquet published to Cloudflare R2.
- Politics section combines open data from the Riksdag, Statistics Sweden, the Election Authority and the Government Offices: votes, speeches, budget proposals and election results, linked to source data.
- Scheduled GitHub Actions jobs refresh the data and publish it; Playwright end-to-end and accessibility tests cover the site, pytest covers the platform.

Tech: React, TypeScript, Vite, Python, dbt, DuckDB, Parquet, Cloudflare R2, GitHub Actions, Playwright

antonernstsson.com **github.com/korv9/anton-portfolio**

Most of my other projects are hosted on the site.

NLP Clustering of ITSM Data

`Spring 2026`

`Graduation thesis · carried out during the Avtalat internship`

- Clustered anonymised incident data from Freshservice and found 72 of 121 recurring incident groups with no link to a registered problem. Several of these were long-running issues worth raising as problem records.
- Measured documentation quality with ISO/IEC 25012 and showed that sparse root-cause notes, rather than the model, were the main limit on the results.
- Compared clustering methods and grounded the findings in interviews with business representatives.

Tech: Python, Azure Databricks, sentence-transformers, UMAP, HDBSCAN, MLflow

DiVA Thesis Lakehouse

`Designed for Microsoft Fabric`

- Incremental OAI-PMH pipeline for thesis metadata with traceable raw XML, Bronze, Silver and Gold layers, Delta MERGE for updated and deleted records, and quality checks before publishing. Tested locally with Spark and Delta Lake.

Tech: PySpark, Delta Lake, Microsoft Fabric, OAI-PMH

Swedish Job Market Analytics

`Personal project`

- Pipeline for JobTech's job ad archive (2020 onwards): Python ingestion and dbt transformations into a star schema, with tests for grain, referential integrity and complete months, CI in GitHub Actions, and semantic clustering of IT job ads.

Tech: dbt Core, DuckDB, SQL, UMAP, HDBSCAN, GitHub Actions

Additional experience

Machine Operator · Delicato AB `2022 - 2025`

Production, quality control and technical troubleshooting. Vice-chair of the local union club.

Languages

Swedish (native) / English (fluent)