

Anton Ernstsson

Junior Data Engineer / Analytics Engineer

Python | Databricks | SQL | Power BI

Stockholm, Sweden | +46 762 541 024 | anton.ernstson@gmail.com

github.com/korv9 | linkedin.com/in/anton-ernstsson

Swedish (native) | English (fluent)

PROFILE

Junior data engineer with experience building production pipelines, from collecting and processing data to making it useful for reporting and analysis. I have worked with stakeholders to understand their needs and improve access to reliable data. I enjoy bringing structure to complex information and understanding how each part of a solution fits together.

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, Spark SQL

Data engineering: Azure Databricks, Delta Lake, dbt Core, DuckDB, ETL/ELT, REST APIs

Architecture and modelling: Lakehouse, Medallion architecture, incremental pipelines, star schema, semantic models

Analytics: Power BI, DAX, DirectQuery, pandas, KPI development

Quality and tooling: Data validation, DQX, pytest, Git, Azure DevOps, GitHub Actions, CI/CD

Applied AI: NLP, sentence-transformers, UMAP, HDBSCAN, MLflow

PROFESSIONAL EXPERIENCE

Analytics Engineer | Avtalat

Jan 2026 - Jun 2026

LIA internship

- Extended a Freshservice ETL pipeline in Azure Databricks using PySpark and Spark SQL, supporting additional instances and analytical use cases.
- Built Power BI reports and semantic models with DirectQuery and DAX, replacing manual reporting with self-service analytics.
- Developed a reproducible NLP analysis of 21,000+ incidents and assessed data quality with attention to GDPR and anonymisation.

Data Engineer | Fora AB

Nov 2025 - Jan 2026

LIA internship

- Built Python, PySpark and Spark SQL pipelines ingesting 80,000+ Freshservice records from a REST API into an Azure Databricks lakehouse.
- Implemented incremental Bronze, Silver and Gold processing with validation, NULL-overwrite protection and error handling.
- Modelled a Gold-layer star schema with hashed surrogate keys for semantic models and ITSM reporting.

EDUCATION

AI Developer | JENSEN Yrkeshögskola

2024 - 2026

Higher Vocational Education (YH), 400 credits, two years full-time. Graduated June 2026.

ADDITIONAL EXPERIENCE

Machine Operator | Delicato AB

2022 - 2025

Production, quality control and technical troubleshooting. Vice-chair of the local union club.

Selected projects

Data platforms, analytics engineering and applied AI

Swedish Job Market Analytics

Analytics engineering pipeline

- Built a pipeline for Swedish job advertisements, with Python handling API and archive ingestion and dbt Core handling SQL transformations, tests and documentation.
- Modelled a star schema with an incremental fact table and tests for grain, referential integrity and calendar completeness; added CI in GitHub Actions.

Python | SQL | dbt Core | DuckDB | Power BI | GitHub Actions

Semantic Drift in Swedish Legislation

Databricks lakehouse and AI evaluation

- Ingested 50 source documents from the Riksdagen API and transformed them into 1,952 structured provisions with immutable snapshots and source traceability.
- Developed a reproducible Python framework to measure meaning shifts across repeated LLM transformations.

Python | PySpark | Delta Lake | Unity Catalog | DuckDB | pytest | REST API

NLP Clustering of ITSM Data

Graduation thesis

- Built an NLP clustering pipeline on 21,000+ production incidents, identifying 72 potential unregistered problem records.
- Evaluated six years of ITSM data against ISO/IEC 25012 and connected data quality and process maturity to model performance.

Python | Databricks | sentence-transformers | UMAP | HDBSCAN | MLflow

Homie API

Python backend and analytics foundation - in progress

- Developing a FastAPI and PostgreSQL service with an explicit OpenAPI contract, authentication, migrations, tests and container configuration.

Python | FastAPI | PostgreSQL | SQLAlchemy | Alembic | Docker Compose